

✓ DELHIVERY BIZ CASE STUDY - FEATURE ENGINEERING

About Delhivery

Delhivery is the largest and fastest-growing fully integrated player in India by revenue in Fiscal 2021. They aim to build the operating system for commerce, through a combination of world-class infrastructure, logistics operations of the highest quality, and cutting-edge engineering and technology capabilities.

The Data team builds intelligence and capabilities using this data that helps them to widen the gap between the quality, efficiency, and profitability of their business versus their competitors.

The company wants to understand and process the data coming out of data engineering pipelines:

- Clean, sanitize and manipulate data to get useful features out of raw fields
- Make sense out of the raw data and help the data science team to build forecasting models on it

✓ **Concept Used:**

1. Feature Creation
2. Relationship between Features
3. Column Normalization /Column Standardization
4. Handling categorical values
5. Missing values - Outlier treatment / Types of outliers

we need functions and methods to do all these analysis, so we must import Python libraries into our work notebook.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as spy
```

To perform Hypothesis testing we need to import few test functions and one hot encoding functions

```
from scipy.stats import ttest_ind, kstest, shapiro
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
```

To get the data into our work space we use the below code(to read csv files) and saving the whole set of data into a single variable(dataframe) which makes analysis easier

```
!wget https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/551/original/delhivery_data.csv?1642751181 -O delh
```

```
--2024-01-12 12:51:22-- https://d2beiqkhq929f0.cloudfront.net/public\_assets/assets/000/001/551/original/delhivery\_data
Resolving d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)... 18.155.174.85, 18.155.174.48, 18.155.174.166
Connecting to d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)|18.155.174.85|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 55617130 (53M) [text/plain]
Saving to: 'delhivery.csv'
```

```
delhivery.csv      100%[=====>]  53.04M  183MB/s   in 0.3s
```

```
2024-01-12 12:51:22 (183 MB/s) - 'delhivery.csv' saved [55617130/55617130]
```

```
df = pd.read_csv('delhivery.csv')  
df.head()
```

```
df.sample()
```

✓ TO ANALYSE THE BASIC METRICS/ BASIC STRUCTURE OF THE DATA:

```
# TO GET NO. OF ROWS & COLUMNS:
```

```
df.shape
```

```
(144867, 24)
```

```
# TO GET TOTAL ELEMENTS IN THE DATASET (i.e., the dot product of no. of rows & columns)
```

```
df.size
```

```
3476808
```

```
# To get index
```

```
df.index
```

```
RangeIndex(start=0, stop=144867, step=1)
```

```
# TO GET THE COLUMNS NAMES
```

```
df.columns
```

```
Index(['data', 'trip_creation_time', 'route_schedule_uuid', 'route_type',  
      'trip_uuid', 'source_center', 'source_name', 'destination_center',  
      'destination_name', 'od_start_time', 'od_end_time',  
      'start_scan_to_end_scan', 'is_cutoff', 'cutoff_factor',  
      'cutoff_timestamp', 'actual_distance_to_destination', 'actual_time',  
      'osrm_time', 'osrm_distance', 'factor', 'segment_actual_time',  
      'segment_osrm_time', 'segment_osrm_distance', 'segment_factor'],  
      dtype='object')
```

```
# TO GET THE NAMES OF THE COLUMNS(alternate method)
```

```
df.keys()
```

```
Index(['data', 'trip_creation_time', 'route_schedule_uuid', 'route_type',
      'trip_uuid', 'source_center', 'source_name', 'destination_center',
      'destination_name', 'od_start_time', 'od_end_time',
      'start_scan_to_end_scan', 'is_cutoff', 'cutoff_factor',
      'cutoff_timestamp', 'actual_distance_to_destination', 'actual_time',
      'osrm_time', 'osrm_distance', 'factor', 'segment_actual_time',
      'segment_osrm_time', 'segment_osrm_distance', 'segment_factor'],
      dtype='object')
```

```
# To get memory usage of each column
```

```
df.memory_usage()
```

```
Index          128
data          1158936
trip_creation_time  1158936
route_schedule_uuid  1158936
route_type     1158936
trip_uuid      1158936
source_center   1158936
source_name     1158936
destination_center  1158936
destination_name  1158936
od_start_time   1158936
od_end_time     1158936
start_scan_to_end_scan  1158936
is_cutoff       144867
cutoff_factor   1158936
cutoff_timestamp  1158936
actual_distance_to_destination  1158936
actual_time     1158936
osrm_time       1158936
osrm_distance   1158936
factor          1158936
segment_actual_time  1158936
segment_osrm_time  1158936
segment_osrm_distance  1158936
```

```
segment_factor          1158936
dtype: int64
```

to get number of unique values in each column

```
df.nunique()
```

```
data                2
trip_creation_time  14817
route_schedule_uuid 1504
route_type           2
trip_uuid           14817
source_center        1508
source_name          1498
destination_center   1481
destination_name     1468
od_start_time        26369
od_end_time          26369
start_scan_to_end_scan 1915
is_cutoff            2
cutoff_factor        501
cutoff_timestamp     93180
actual_distance_to_destination 144515
actual_time          3182
osrm_time            1531
osrm_distance        138046
factor              45641
segment_actual_time   747
segment_osrm_time     214
segment_osrm_distance 113799
segment_factor        5675
dtype: int64
```

```
# To get the Time period for which the data is been taken
```

```
mini = df['trip_creation_time'].min()
maxi = df['od_end_time'].max()
print(f'start period : {mini}')
print(f'end period : {maxi}')

start period : 2018-09-12 00:00:16.535741
end period : 2018-10-08 03:00:24.353479
```

✓ INFERENCE:

The given data is form the year **2018** and confined from **12th september** and **08th october** months.

```
# TO GET THE STATISTICAL SUMMARY:
```

```
df.describe().T
```

```
df.describe(include = object).T
```



```
df['data'].value_counts()
```

```
training    104858
test         40009
Name: data, dtype: int64
```

```
df['route_type'].value_counts()
```

```
FTL          99660
Carting       45207
Name: route_type, dtype: int64
```

```
df['source_name'].value_counts()
```

```

Gurgaon_Bilaspur_HB (Haryana)      23347
Bangalore_Nelmngla_H (Karnataka)    9975
Bhiwandi_Mankoli_HB (Maharashtra)   9088
Pune_Tathawde_H (Maharashtra)       4061
Hyderabad_Shamshbd_H (Telangana)     3340
...
Shahjhnpur_NavdaCln_D (Uttar Pradesh) 1
Soro_UttarDPP_D (Orissa)             1
Kayamkulam_Bhrnikvu_D (Kerala)       1
Krishnanagar_AnadiDPP_D (West Bengal) 1
Faridabad_Old (Haryana)              1
Name: source_name, Length: 1498, dtype: int64

```

```
df['destination_name'].value_counts()
```

```

Gurgaon_Bilaspur_HB (Haryana)      15192
Bangalore_Nelmngla_H (Karnataka)    11019
Bhiwandi_Mankoli_HB (Maharashtra)   5492
Hyderabad_Shamshbd_H (Telangana)     5142
Kolkata_Dankuni_HB (West Bengal)     4892
...
Hyd_Trimulgherry_Dc (Telangana)      1
Vijayawada (Andhra Pradesh)          1
Baghpat_Barout_D (Uttar Pradesh)     1
Mumbai_Sanpada_CP (Maharashtra)      1
Basta_Central_DPP_1 (Orissa)         1
Name: destination_name, Length: 1468, dtype: int64

```

INFERENCE FROM RAW DATA:

1. There are more **training**(104858) data
2. **FTL**(Full Truck Load)(99660) - is the most preferred transportation type
3. Most trip is originated at **Gurgaon_Bilaspur_HB (Haryana)** (23347)
4. Majority of the trips are destined to **Gurgaon_Bilaspur_HB (Haryana)** (15192)

✓ DROPPING ALL THE UNKNOWN FIELDS

The unknown fields has been removed from the original data to perform further analysis.

```
unknown_fields = ['is_cutoff', 'cutoff_factor', 'cutoff_timestamp', 'factor', 'segment_factor']
df.drop(columns = unknown_fields, inplace=True)
```

```
df.columns
```

```
Index(['data', 'trip_creation_time', 'route_schedule_uuid', 'route_type',
      'trip_uuid', 'source_center', 'source_name', 'destination_center',
      'destination_name', 'od_start_time', 'od_end_time',
      'start_scan_to_end_scan', 'actual_distance_to_destination',
      'actual_time', 'osrm_time', 'osrm_distance', 'segment_actual_time',
      'segment_osrm_time', 'segment_osrm_distance'],
      dtype='object')
```

✓ MISSING VALUE DETECTION

```
df.isnull().sum()
```

```
data          0
trip_creation_time  0
route_schedule_uuid  0
route_type    0
trip_uuid     0
source_center  0
source_name   293
destination_center  0
destination_name  261
od_start_time  0
od_end_time    0
start_scan_to_end_scan  0
actual_distance_to_destination  0
actual_time    0
osrm_time      0
```

```
osrm_distance          0
segment_actual_time    0
segment_osrm_time      0
segment_osrm_distance  0
dtype: int64
```

INFERENCE:

There are no null values/missing values in the dataframe except for two columns,

1. source_name - which has 293 missing values
2. destination_name - which has 261 missing values

✓ NUMBER OF UNKNOWN SOURCE & DESTINATION CENTERS:

```
df.loc[df['source_name'].isnull()]['source_center'].nunique()
```

```
10
```

```
df.loc[df['destination_name'].isnull()]['destination_center'].nunique()
```

```
13
```

✓ DEALING WITH NULL VALUES:

```
df['source_name'].fillna('No name',inplace=True)
df['destination_name'].fillna('No name',inplace=True)
```

```
df.isnull().sum()
```

```

data                0
trip_creation_time  0
route_schedule_uuid 0
route_type          0
trip_uuid           0
source_center       0
source_name         0
destination_center  0
destination_name     0
od_start_time       0
od_end_time         0
start_scan_to_end_scan 0
actual_distance_to_destination 0
actual_time         0
osrm_time           0
osrm_distance       0
segment_actual_time 0
segment_osrm_time   0
segment_osrm_distance 0
dtype: int64

```

✓ INFERENCE:

The null values have been replaced by 'No name'.

```
df.nunique()
```

```

data                2
trip_creation_time  14817
route_schedule_uuid 1504
route_type          2
trip_uuid           14817
source_center       1508
source_name         1499
destination_center  1481
destination_name     1469
od_start_time       26369
od_end_time         26369
start_scan_to_end_scan 1915

```

```

actual_distance_to_destination    144515
actual_time                      3182
osrm_time                       1531
osrm_distance                   138046
segment_actual_time              747
segment_osrm_time                214
segment_osrm_distance            113799
dtype: int64

```

```
# TO GET THE TOTAL INFORMATION ABOUT THE DATASET.
```

```
# info function let us know the columns with their data types and no. of non-null values & the total memory usage
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   data                                  144867 non-null  object
1   trip_creation_time                  144867 non-null  object
2   route_schedule_uuid                144867 non-null  object
3   route_type                         144867 non-null  object
4   trip_uuid                          144867 non-null  object
5   source_center                      144867 non-null  object
6   source_name                        144867 non-null  object
7   destination_center                 144867 non-null  object
8   destination_name                   144867 non-null  object
9   od_start_time                      144867 non-null  object
10  od_end_time                        144867 non-null  object
11  start_scan_to_end_scan              144867 non-null  float64
12  actual_distance_to_destination      144867 non-null  float64
13  actual_time                        144867 non-null  float64
14  osrm_time                          144867 non-null  float64
15  osrm_distance                      144867 non-null  float64
16  segment_actual_time                 144867 non-null  float64
17  segment_osrm_time                  144867 non-null  float64
18  segment_osrm_distance               144867 non-null  float64
dtypes: float64(8), object(11)
memory usage: 21.0+ MB

```

INFERENCE:

We can see that,

1. Columns involving time are in different data type
2. Also, columns with 2 unique entries implies that they are categorical

✓ CHANGING THE DTYPE:

```
df[['data', 'route_type']] = df[['data', 'route_type']].astype('category')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   data                                  144867 non-null  category
1   trip_creation_time                   144867 non-null  object
2   route_schedule_uuid                 144867 non-null  object
3   route_type                           144867 non-null  category
4   trip_uuid                           144867 non-null  object
5   source_center                       144867 non-null  object
6   source_name                         144867 non-null  object
7   destination_center                  144867 non-null  object
8   destination_name                    144867 non-null  object
9   od_start_time                      144867 non-null  object
10  od_end_time                         144867 non-null  object
11  start_scan_to_end_scan              144867 non-null  float64
12  actual_distance_to_destination      144867 non-null  float64
13  actual_time                         144867 non-null  float64
14  osrm_time                          144867 non-null  float64
15  osrm_distance                      144867 non-null  float64
```

```

16 segment_actual_time          144867 non-null float64
17 segment_osrm_time            144867 non-null float64
18 segment_osrm_distance        144867 non-null float64
dtypes: category(2), float64(8), object(9)
memory usage: 19.1+ MB

```

✓ TO CHANGE INTO DATETIME DTYPE:

```

datetime_columns = ['trip_creation_time', 'od_start_time', 'od_end_time']
for element in datetime_columns:
    df[element] = pd.to_datetime(df[element])

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26368 entries, 0 to 26367
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   trip_uuid                            26368 non-null  object
1   source_center                        26368 non-null  object
2   destination_center                   26368 non-null  object
3   data                                 26368 non-null  category
4   route_type                           26368 non-null  category
5   trip_creation_time                   26368 non-null  datetime64[ns]
6   source_name                          26368 non-null  object
7   destination_name                     26368 non-null  object
8   od_start_time                       26368 non-null  datetime64[ns]
9   od_end_time                         26368 non-null  datetime64[ns]
10  start_scan_to_end_scan               26368 non-null  float64
11  actual_distance_to_destination        26368 non-null  float64
12  actual_time                           26368 non-null  float64
13  osrm_time                            26368 non-null  float64
14  osrm_distance                        26368 non-null  float64
15  segment_actual_time                   26368 non-null  float64
16  segment_osrm_time                     26368 non-null  float64
17  segment_osrm_distance                 26368 non-null  float64

```



```
dtypes: category(2), datetime64[ns](3), float64(8), object(5)  
memory usage: 3.3+ MB
```

INFERENCE:

Now, We can clearly visualise the change in dtype of few columns which are more relevant to do analysis

✓ TO MERGE THE ROWS AND AGGREGATING FIELDS

Hint: We can use inbuilt functions like groupby and aggregations like sum(), cumsum() to merge some rows based on their,

1. Trip_uuid, Source ID and Destination ID
2. Further aggregate on the basis of just Trip_uuid. We can also keep the first and last values for some numeric/categorical fields if aggregating them won't make sense.

```
df = df.groupby(by = ['trip_uuid', 'source_center', 'destination_center'], as_index=False).agg({'data' : 'first',  
                                                'route_type' : 'first',  
                                                'trip_creation_time' : 'first',  
                                                'source_name' : 'first',  
                                                'destination_name' : 'last',  
                                                'od_start_time' : 'first',  
                                                'od_end_time' : 'first',  
                                                'start_scan_to_end_scan' : 'first',  
                                                'actual_distance_to_destination' : 'last',  
                                                'actual_time' : 'last',  
                                                'osrm_time' : 'last',  
                                                'osrm_distance' : 'last',  
                                                'segment_actual_time' : 'sum',  
                                                'segment_osrm_time' : 'sum',  
                                                'segment_osrm_distance' : 'sum'})
```

```
df.sample(5)
```

```
df['trip_uuid'].nunique()
```

```
14817
```

- ✓ Calculate the time taken between od_start_time and od_end_time and keep it as a feature.
- ✓ Drop the original columns, if required

```
df['od_total_time'] = df['od_end_time'] - df['od_start_time']  
df.drop(columns = ['od_end_time', 'od_start_time'], inplace = True)  
df['od_total_time'] = df['od_total_time'].apply(lambda x : round(x.total_seconds() / 60.0, 2))  
df['od_total_time'].head()
```

```
0    1260.60  
1     999.51  
2      58.83  
3     122.78
```

```
4      834.64
```

```
Name: od_total_time, dtype: float64
```

```
df = df.groupby(by = 'trip_uuid', as_index = False).agg({'source_center' : 'first',  
                                                         'destination_center' : 'last',  
                                                         'data' : 'first',  
                                                         'route_type' : 'first',  
                                                         'trip_creation_time' : 'first',  
                                                         'source_name' : 'first',  
                                                         'destination_name' : 'last',  
                                                         'od_total_time' : 'sum',  
                                                         'start_scan_to_end_scan' : 'sum',  
                                                         'actual_distance_to_destination' : 'sum',  
                                                         'actual_time' : 'sum',  
                                                         'osrm_time' : 'sum',  
                                                         'osrm_distance' : 'sum',  
                                                         'segment_actual_time' : 'sum',  
                                                         'segment_osrm_time' : 'sum',  
                                                         'segment_osrm_distance' : 'sum'})  
  
df.head()
```

- ✓ Build some features to prepare the data for actual analysis. Extract features from the below fields:

```
df['source_name'].head(5)
```

```
0    Anand_VUNagar_DC (Gujarat)
1    Anand_VUNagar_DC (Gujarat)
2    Anand_VUNagar_DC (Gujarat)
3    Anand_VUNagar_DC (Gujarat)
4    Anand_VUNagar_DC (Gujarat)
Name: source_name, dtype: object
```

```
def extract_city(x):  
    if x == 'No name':  
        return 'unknown_city'  
    else:  
        l = x.split()[0].split('_')  
        if 'CCU' in x:  
            return 'Kolkata'  
        elif 'MAA' in x.upper():  
            return 'Chennai'  
        elif ('HBR' in x.upper()) or ('BLR' in x.upper()):  
            return 'Bengaluru'  
        elif 'FBD' in x.upper():  
            return 'Faridabad'  
        elif 'BOM' in x.upper():  
            return 'Mumbai'  
        elif 'DEL' in x.upper():  
            return 'Delhi'  
        elif 'OK' in x.upper():  
            return 'Delhi'  
        elif 'GZB' in x.upper():  
            return 'Ghaziabad'  
        elif 'GGN' in x.upper():  
            return 'Gurgaon'  
        elif 'AMD' in x.upper():  
            return 'Ahmedabad'  
        elif 'CJB' in x.upper():  
            return 'Coimbatore'  
        elif 'HYD' in x.upper():  
            return 'Hyderabad'  
        return l[0]
```

```
def extract_state(x):
    if x == 'No name':
        return 'State not found'
    else:
        temp = x.split('(')
        if len(temp) == 1:
            return temp[0]
        else:
            return temp[1].replace(')','')
```

```
def extract_place(x):
    if 'No name' in x:
        return x
    elif 'HBR' in x:
        return 'HBR Layout PC'
    else:
        l = x.split()[0].split('_', 1)
        if len(l) == 1:
            return 'unknown_place'
        else:
            return l[1]
```

✓ 1. Source Name: Split and extract features out of destination. City-place-code (State)

```
df['source_state'] = df['source_name'].apply(extract_state)

print('Number of source states : ',df['source_state'].nunique())

df['source_state'].unique()
```

```
Number of source states : 30
array(['Uttar Pradesh', 'Karnataka', 'Haryana', 'Maharashtra',
       'Tamil Nadu', 'Gujarat', 'Delhi', 'Telangana', 'Rajasthan',
       'Assam', 'Madhya Pradesh', 'West Bengal', 'Andhra Pradesh',
       'Punjab', 'Chandigarh', 'Goa', 'Jharkhand', 'Pondicherry',
```

```
'Orissa', 'Uttarakhand', 'Himachal Pradesh', 'Kerala',
'Arunachal Pradesh', 'Bihar', 'Chhattisgarh',
'Dadra and Nagar Haveli', 'Jammu & Kashmir', 'Mizoram', 'Nagaland',
'State not found'], dtype=object)
```

```
df['source_city'] = df['source_name'].apply(extract_city)
print('No of source cities :', df['source_city'].nunique())
df['source_city'].unique()[:50]
```

No of source cities : 690

```
array(['Kanpur', 'Doddablpur', 'Gurgaon', 'Mumbai', 'Bellary', 'Chennai',
'Bengaluru', 'Surat', 'Delhi', 'Pune', 'Faridabad', 'Shirala',
'Hyderabad', 'Thirumalagiri', 'Gulbarga', 'Jaipur', 'Allahabad',
'Guwahati', 'Narsinghpur', 'Shrirampur', 'Madakasira', 'Sonari',
'Dindigul', 'Jalandhar', 'Chandigarh', 'Deoli', 'Pandharpur',
'Kolkata', 'Bhandara', 'Kurnool', 'Bhiwandi', 'Bhatinda',
'RoopNagar', 'Bantwal', 'Lalru', 'Kadi', 'Shahdol', 'Gangakher',
'Durgapur', 'Vapi', 'Jamjodhpur', 'Jetpur', 'Mehsana', 'Jabalpur',
'Junagadh', 'Gundlupet', 'Mysore', 'Goa', 'Bhopal', 'Sonipat'],
dtype=object)
```

```
df['source_place'] = df['source_name'].apply(extract_place)
print('Number of source places : ', df['source_place'].nunique())
df['source_place'].unique()[:50]
```

Number of source places : 757

```
array(['Central_H_6', 'ChikaDPP_D', 'Bilaspur_HB', 'unknown_place', 'Dc',
'Poonamallee', 'Chrompet_DPC', 'HBR Layout PC', 'Central_D_12',
'Lajpat_IP', 'North_D_3', 'Balabgarh_DPC', 'Central_DPP_3',
'Shamshbd_H', 'Xroad_D', 'Nehrugn_I', 'Central_I_7',
'Central_H_1', 'Nangli_IP', 'North', 'KndliDPP_D', 'Central_D_9',
'DavkharRd_D', 'Bandel_D', 'RTCStand_D', 'Central_DPP_1',
'KGAirprt_HB', 'North_D_2', 'Central_D_1', 'DC', 'Mthurard_L',
'Mullanpr_DC', 'Central_DPP_2', 'RajCmplx_D', 'Beliaghata_DPC',
'RjnaiDPP_D', 'AbbasNgr_I', 'Mankoli_HB', 'DPC', 'Airport_H',
'Hub', 'Gateway_HB', 'Tathawde_H', 'ChotiHvl_DC', 'Trmltpl_D',
'OnkarDPP_D', 'Mehmdpur_H', 'KaranNGR_D', 'Sohagpur_D',
'Chrompet_L'], dtype=object)
```

✓ 2. Destination Name: Split and extract features out of destination. City-place-code (State)

```
df['destination_state'] = df['destination_name'].apply(extract_state)
print('Number of destination states : ',df['destination_state'].nunique())
df['destination_state'].unique()
```

Number of destination states : 32

```
array(['Uttar Pradesh', 'Karnataka', 'Haryana', 'Maharashtra',
      'Tamil Nadu', 'Gujarat', 'Delhi', 'Telangana', 'Rajasthan',
      'Madhya Pradesh', 'Assam', 'West Bengal', 'Andhra Pradesh',
      'Punjab', 'Chandigarh', 'Dadra and Nagar Haveli', 'Orissa',
      'Bihar', 'Jharkhand', 'Goa', 'Uttarakhand', 'Himachal Pradesh',
      'Kerala', 'Arunachal Pradesh', 'Mizoram', 'Chhattisgarh',
      'Jammu & Kashmir', 'Nagaland', 'Meghalaya', 'Tripura',
      'State not found', 'Daman & Diu'], dtype=object)
```

```
df['destination_city'] = df['destination_name'].apply(extract_city)
print('Number of destination city : ',df['destination_city'].nunique())
df['destination_city'].unique()[:50]
```

Number of destination city : 806

```
array(['Kanpur', 'Doddablpur', 'Gurgaon', 'Mumbai', 'Sandur', 'Chennai',
      'Bengaluru', 'Surat', 'Delhi', 'PNQ', 'Faridabad', 'Ratnagiri',
      'Bangalore', 'Hyderabad', 'Aland', 'Jaipur', 'Satna', 'Guwahati',
      'Bareli', 'Nashik', 'Hooghly', 'Sivasagar', 'Palani', 'Jalandhar',
      'Chandigarh', 'Yavatmal', 'Sangola', 'Kolkata', 'Savner',
      'Kurnool', 'Bhatinda', 'Bhiwandi', 'Barnala', 'Murbad', 'Kadaba',
      'Gulbarga', 'Naraingarh', 'Ludhiana', 'Kadi', 'Jabalpur',
      'Gangakher', 'Bankura', 'Silvassa', 'Porbandar', 'Jetpur',
      'Khammam', 'Mehsana', 'Katni', 'Una', 'Malavalli'], dtype=object)
```



```
df['destination_place'] = df['destination_name'].apply(extract_place)
print('Number of destination place : ',df['destination_place'].nunique())
df['destination_place'].unique()[:50]
```

```
Number of destination place : 843
array(['Central_H_6', 'ChikaDPP_D', 'Bilaspur_HB', 'MiraRd_IP',
      'WrdN1DPP_D', 'Poonamallee', 'Vandalur_Dc', 'HBR Layout PC',
      'Central_D_3', 'Bhogal', 'unknown_place', 'MjgaonRd_D',
      'Nelmngla_H', 'Uppal_I', 'RazaviRd_D', 'Central_I_7',
      'Central_I_2', 'Hub', 'SourvDPP_D', 'Varachha_DC', 'TgrniaRD_I',
      'DC', 'Gokulam_D', 'Babupaty_D', 'Bomsndra_HB', 'Alwal_I',
      'RjndraRd_D', 'Mehmdpur_H', 'Sanpada_I', 'JajuDPP_D',
      'Central_DPP_2', 'Dankuni_HB', 'Wagodha_D', 'AbbasNgr_I',
      'Balabgharh_DPC', 'DPC', 'Mankoli_HB', 'Shamshbd_H', 'SnkunDPP_D',
      'Kharar_DC', 'AnugrDPP_D', 'Nehrugn_I', 'Ward2DPP_D',
      'MilrGanj_HB', 'KaranNGR_D', 'Adhartal_IP', 'Poonamallee_HB',
      'Busstand_D', 'BhowmDPP_D', 'Samrvrni_D'], dtype=object)
```

✓ 3. Trip_creation_time: Extract features like month, year and day etc

```
df['trip_creation_time'].head(2)
```

```
0    2018-09-12 00:00:16.535741
1    2018-09-12 00:00:22.886430
Name: trip_creation_time, dtype: datetime64[ns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14817 entries, 0 to 14816
Data columns (total 23 columns):
#   Column              Non-Null Count  Dtype
---  -
0   trip_uuid           14817 non-null  object
1   source_center       14817 non-null  object
2   destination_center  14817 non-null  object
3   data                14817 non-null  category
```

```

4  route_type                14817 non-null category
5  trip_creation_time        14817 non-null datetime64[ns]
6  source_name               14817 non-null object
7  destination_name          14817 non-null object
8  od_total_time             14817 non-null float64
9  start_scan_to_end_scan    14817 non-null float64
10 actual_distance_to_destination 14817 non-null float64
11 actual_time               14817 non-null float64
12 osrm_time                 14817 non-null float64
13 osrm_distance             14817 non-null float64
14 segment_actual_time       14817 non-null float64
15 segment_osrm_time         14817 non-null float64
16 segment_osrm_distance     14817 non-null float64
17 source_state              14817 non-null object
18 source_city               14817 non-null object
19 source_place              14817 non-null object
20 destination_state         14817 non-null object
21 destination_city          14817 non-null object
22 destination_place         14817 non-null object
dtypes: category(2), datetime64[ns](1), float64(9), object(11)
memory usage: 2.4+ MB

```

```

df['trip_creation_year'] = df['trip_creation_time'].dt.year
df['trip_creation_month'] = df['trip_creation_time'].dt.month
df['trip_creation_day'] = df['trip_creation_time'].dt.day
df['trip_creation_week'] = df['trip_creation_time'].dt.isocalendar().week
df['trip_creation_week'] = df['trip_creation_week'].astype('int8')

```

```

df['trip_creation_hour'] = df['trip_creation_time'].dt.hour
df['trip_creation_time'] = df['trip_creation_time'].dt.strftime('%H:%M:%S')

```

✓ Data after cleaning:

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14817 entries, 0 to 14816
Data columns (total 28 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   trip_uuid                            14817 non-null  object
 1   source_center                        14817 non-null  object
 2   destination_center                   14817 non-null  object
 3   data                                 14817 non-null  category
 4   route_type                           14817 non-null  category
 5   trip_creation_time                   14817 non-null  object
 6   source_name                          14817 non-null  object
 7   destination_name                     14817 non-null  object
 8   od_total_time                        14817 non-null  float64
 9   start_scan_to_end_scan               14817 non-null  float64
10   actual_distance_to_destination        14817 non-null  float64
11   actual_time                          14817 non-null  float64
12   osrm_time                            14817 non-null  float64
13   osrm_distance                        14817 non-null  float64
14   segment_actual_time                  14817 non-null  float64
15   segment_osrm_time                    14817 non-null  float64
16   segment_osrm_distance                 14817 non-null  float64
17   source_state                         14817 non-null  object
18   source_city                          14817 non-null  object
19   source_place                         14817 non-null  object
20   destination_state                    14817 non-null  object
21   destination_city                     14817 non-null  object
22   destination_place                    14817 non-null  object
23   trip_creation_year                   14817 non-null  int64
24   trip_creation_month                   14817 non-null  int64
25   trip_creation_day                     14817 non-null  int64
26   trip_creation_week                    14817 non-null  int8
27   trip_creation_hour                    14817 non-null  int64
dtypes: category(2), float64(9), int64(4), int8(1), object(12)
memory usage: 2.9+ MB

```

```
df.shape
```

```
(14817, 28)
```

```
df.describe().T
```

```
df.describe(include = object).T
```

TO GET MAX AND MIN PERIOD OF THE TRIP CREATION(ANALYSING MONTHS,WEEK, DAY, HOUR)

✓ ANALYSIS BASED ON DAYS OF THE MONTH

```
df['trip_creation_day'].unique()
```

```
array([12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28,  
       29, 30,  1,  2,  3])
```

```
df_day = df.groupby(by = 'trip_creation_day')['trip_uuid'].count().reset_index()  
df_day.sample(3)
```

```
plt.figure(figsize = (12, 4))
sns.lineplot(data = df_day,
             x = df_day['trip_creation_day'],
             y = df_day['trip_uuid'])
plt.xticks(np.arange(1, 32))
plt.title('DAY-WISE ANALYSIS')
plt.grid()
```

INFERENCE:

Most trips are created in the **middle** of the month.

✓ ANALYSIS BASED ON WEEK

```
df['trip_creation_week'].unique()
```

```
array([37, 38, 39, 40], dtype=int8)
```

```
df_week = df.groupby(by = 'trip_creation_week')['trip_uuid'].count().to_frame().reset_index()  
df_week.head()
```

```
plt.figure(figsize = (12, 3))  
sns.lineplot(data = df_week,  
             x = df_week['trip_creation_week'],  
             y = df_week['trip_uuid'])  
plt.title('WEEKLY ANALYSIS')  
plt.grid()
```

INFERENCE:

Maximum trips are created in the **38th** week

✓ ANALYSIS BASED ON THE HOURS OF THE DAY

```
df['trip_creation_hour'].unique()
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,  
       17, 18, 19, 20, 21, 22, 23])
```

```
df_hour = df.groupby(by = 'trip_creation_hour')['trip_uuid'].count().reset_index()  
df_hour.head()
```



```
plt.figure(figsize = (12, 4))
sns.lineplot(data = df_hour,
             x = df_hour['trip_creation_hour'],
             y = df_hour['trip_uuid'])
plt.xticks(np.arange(0,24))
plt.title('HOURLY ANALYSIS')
plt.grid()
```

INFERENCE:

1. During early hours of the day, the count is good
2. But after 5AM till 12PM the count is reducing
3. After 1pm the count starts gradually increasing
4. Reaches **maximum** count on **10PM**

✓ ANALYSIS BASED ON TYPE OF DATA

```
df_data = df.groupby('data')['trip_uuid'].count().reset_index()
df_data['percentage'] = np.round(df_data['trip_uuid'] * 100/ df_data['trip_uuid'].sum(), 2)
df_data.head()
```

```
plt.figure(figsize=(3,3))
plt.pie(x = df_data['trip_uuid'],
        labels = df_data['data'],
        explode = [0, 0.1],
        autopct = '%.2f%%')
plt.title('TEST VS TRAINING')
plt.plot()
```

INFERENCE:

Training data type has maximum trips

✓ ANALYSIS BASED ON ROUTE TYPE

```
df_route = df.groupby(by = 'route_type')['trip_uuid'].count().to_frame().reset_index()
df_route['percentage'] = np.round(df_route['trip_uuid'] * 100/ df_route['trip_uuid'].sum(), 2)
df_route.head()
```

```
plt.figure(figsize=(3,3))
plt.pie(x = df_route['trip_uuid'],
        labels = ['Carting', 'FTL'],
        explode = [0, 0.1],
        autopct = '%.2f%%')
plt.title('ROUTE TYPE ANALYSIS')
plt.plot()
```

INFERENCE:

Carting route type has maximum trips

✓ DISTRIBUTION OF NUMBER OF TRIPS CREATED IN DIFFERENT STATES

```
df.describe(include = object).T
```

```
df_source_state = df.groupby('source_state')['trip_uuid'].count().reset_index()
df_source_state['percentage'] = np.round(df_source_state['trip_uuid'] * 100 / df_source_state['trip_uuid'].sum(), 2)
df_source_state = df_source_state.sort_values(by = 'trip_uuid', ascending = False)
df_source_state.head(2)
```

```
plt.figure(figsize = (10, 15))
sns.barplot(data = df_source_state,
            x = df_source_state['trip_uuid'],
            y = df_source_state['source_state'])
plt.title('Source state - Analysis')
plt.show()
```


INFERENCE:

Maximum trips are originated from **Maharashtra** followed by karnataka and Haryana

✓ Distribution of trips created in different cities

```
df_source_city = df.groupby('source_city')['trip_uuid'].count().reset_index()
df_source_city['percentage'] = np.round(df_source_city['trip_uuid'] * 100/ df_source_city['trip_uuid'].sum(), 2)
df_source_city = df_source_city.sort_values(by = 'trip_uuid', ascending = False)[:30]
df_source_city.head(2)
```

```
plt.figure(figsize = (10, 10))
sns.barplot(data = df_source_city,
            x = df_source_city['trip_uuid'],
            y = df_source_city['source_city'])
plt.title('Trips oriented in different cities - analysis')
plt.show()
```


INFERENCE:

Most trips are originated in **Mumbai** followed by Gurgaon and Delhi

✓ Analysis based on the distribution of number of trips based on destination state

```
df_destination_state = df.groupby('destination_state')['trip_uuid'].count().reset_index()
df_destination_state['percentage'] = np.round(df_destination_state['trip_uuid'] * 100/ df_destination_state['trip_uuid'].su
df_destination_state = df_destination_state.sort_values(by = 'trip_uuid', ascending = False)
df_destination_state.head(2)
```

```
plt.figure(figsize = (10, 10))
sns.barplot(data = df_destination_state,
            x = df_destination_state['trip_uuid'],
            y = df_destination_state['destination_state'])
plt.title('Destination state - Analysis')
plt.show()
```


INFERENCE:

Maximum trips are destined to **Maharastra**

✓ Analysis based on the distribution of number of trips based on destination cities

```
df_destination_city = df.groupby('destination_city')['trip_uuid'].count().reset_index()
df_destination_city['percentage'] = np.round(df_destination_city['trip_uuid'] * 100/ df_destination_city['trip_uuid'].sum())
df_destination_city = df_destination_city.sort_values(by = 'trip_uuid', ascending = False)[:30]
df_destination_city.head(3)
```

```
plt.figure(figsize = (10, 10))
sns.barplot(data = df_destination_city,
            x = df_destination_city['trip_uuid'],
            y = df_destination_city['destination_city'])
plt.title('Destination city - Analysis')
plt.show()
```


INFERENCE:

maximum trips are destined to **Mumbai** city followed by Bengaluru and Gurgaon

✓ PAIRPLOT:

```
numerical_columns = ['od_total_time', 'start_scan_to_end_scan', 'actual_distance_to_destination',  
                    'actual_time', 'osrm_time', 'osrm_distance', 'segment_actual_time',  
                    'segment_osrm_time', 'segment_osrm_distance']  
sns.pairplot(data = df,  
            vars = numerical_columns,  
            kind = 'reg',  
            hue = 'route_type',  
            markers = '*')  
plt.show()
```



```
df_corr = df[numerical_columns].corr()  
df_corr
```

```
plt.figure(figsize = (15, 10))  
sns.heatmap(data = df_corr, annot = True)  
plt.show()
```


INFERENCE:

Very High Correlation (> 0.9) exists between columns all the numerical columns specified above

✓ 3. In-depth analysis and feature engineering:

- ✓ Compare the difference between `od_total_time` and `start_scan_to_end_scan`. Do hypothesis testing/ Visual analysis to check.

STEP-1 : Set up Null Hypothesis

1. **Null Hypothesis (H_0)** - `od_total_time` (Total Trip Time) and `start_scan_to_end_scan` (Expected total trip time) are same.
 2. **Alternate Hypothesis (H_A)** - `od_total_time` (Total Trip Time) and `start_scan_to_end_scan` (Expected total trip time) are different.
-

STEP-2 : Checking for basic assumptions for the hypothesis

Distribution check using **QQ Plot**

Homogeneity of Variances using **Lavene's test**

STEP-3: Define Test statistics; Distribution of T under H_0 .

If the assumptions of T Test are met then we can proceed performing T Test for independent samples else we will perform the non parametric test equivalent to T Test for independent sample i.e., Mann-Whitney U rank test for two independent samples.

STEP-4: Compute the p-value and fix value of alpha.

We set our **alpha to be 0.05**

STEP-5: Compare p-value and alpha.

Based on p-value, we will reject or fail to reject H_0 .

p-val < alpha : Reject H_0

p-val > alpha : Fail to reject H_0

```
df[['od_total_time', 'start_scan_to_end_scan']].describe().T
```

✓ KDE PLOT:

```
plt.figure(figsize=(3,3))
sns.kdeplot(df['od_total_time'])
sns.kdeplot(df['start_scan_to_end_scan'])
plt.show()
```

INFERENCE:

The **kdeplot** vividly shows that the graphs of both the groups are the **same distribution** and they have **almost same mean**. so ttest is performed.

✓ TTEST:

```
# H0 - od_total_time and start_scan_to_end_scan are similar
# Ha - od_total_time and start_scan_to_end_scan are different

test_stat, p_value = spy.ttest_ind(df['od_total_time'],df['start_scan_to_end_scan'])

print('P-value :',p_value)
alpha = 0.05

if p_value < alpha:
    print('REJECT H0 - od_total_time and start_scan_to_end_scan are different')

else:
    print('FAIL TO REJECT H0 - od_total_time and start_scan_to_end_scan are similar')

    P-value : 0.9076773879740099
    FAIL TO REJECT H0 - od_total_time and start_scan_to_end_scan are similar
```

Visual Tests to know if the samples follow normal distribution

✓ QQPLOT:

```
plt.figure(figsize = (15, 5))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for od_total_time and start_scan_to_end_scan')
spy.probplot(df['od_total_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for od_total_time')
plt.subplot(1, 2, 2)
spy.probplot(df['start_scan_to_end_scan'], plot = plt, dist = 'norm')
plt.title('QQ plot for start_scan_to_end_scan')
plt.show()
```


It can be seen from the above plots that the samples do not come from normal distribution.

Applying Shapiro-Wilk test for normality

H_0 : The sample follows normal distribution

H_a : The sample does not follow normal distribution

$\alpha = 0.05$

Test Statistics : Shapiro-Wilk test for normality

✓ SHAPIRO-WILK TEST:

```
test_stat, p_value = shapiro(df['od_total_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT  $H_0$  - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT  $H_0$  - The sample follows normal distribution')
```

```
p-value 0.0  
REJECT H0 - The sample does not follow normal distribution
```

```
test_stat, p_value = shapiro(df['start_scan_to_end_scan'].sample(5000))  
print('p-value', p_value)  
if p_value < 0.05:  
    print('REJECT H0 - The sample does not follow normal distribution')  
else:  
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0  
REJECT H0 - The sample does not follow normal distribution
```

✓ KS test for normality

```
test_stat, p_value = kstest(df['od_total_time'].sample(5000), spy.norm.cdf)  
print('p-value', p_value)  
if p_value < 0.05:  
    print('REJECT H0 - The sample does not follow normal distribution')  
else:  
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0  
REJECT H0 - The sample does not follow normal distribution
```

```
test_stat, p_value = kstest(df['start_scan_to_end_scan'].sample(5000), spy.norm.cdf)  
print('p-value', p_value)  
if p_value < 0.05:  
    print('REJECT H0 - The sample does not follow normal distribution')  
else:  
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0  
REJECT H0 - The sample does not follow normal distribution
```

Transforming the data using boxcox transformation to check if the transformed data follows normal distribution.

✓ BOXCOX TRANSFORMATION:

```
transformed_od_total_time = spy.boxcox(df['od_total_time'])[0]
test_stat, p_value = spy.shapiro(transformed_od_total_time)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 7.172770042757021e-25
REJECT H0 -The sample does not follow normal distribution
/usr/local/lib/python3.10/dist-packages/scipy/stats/_morestats.py:1882: UserWarning: p-value may not be accurate for N
    warnings.warn("p-value may not be accurate for N > 5000.")
```

```
transformed_start_scan_to_end_scan = spy.boxcox(df['start_scan_to_end_scan'])[0]
test_stat, p_value = spy.shapiro(transformed_start_scan_to_end_scan)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 1.0471322892609475e-24
REJECT H0 -The sample does not follow normal distribution
```

For all the test, the sample doesnt follow Gaussian distribution

Homogeneity of Variances using Lavene's test

✓ LAVENE'S TEST:

```
# H0 - Homogenous Variance

# HA - Non Homogenous Variance

test_stat, p_value = spy.levene(df['od_total_time'], df['start_scan_to_end_scan'])
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The samples do not have Homogenous Variance')
else:
    print('FAIL TO REJECT H0 - The samples have Homogenous Variance ')

p-value 0.9668007217581142
FAIL TO REJECT H0 - The samples have Homogenous Variance
```

Since the samples are not normally distributed, T-Test cannot be applied here, we can perform its non parametric equivalent test i.e., **Mann-Whitney U rank** test for two independent samples.

✓ Mann-Whitney U rank test:

```
# H0 - od_total_time and start_scan_to_end_scan are similar
# Ha - od_total_time and start_scan_to_end_scan are different

test_stat, p_value = spy.mannwhitneyu(df['od_total_time'], df['start_scan_to_end_scan'])
print('P-value :',p_value)
alpha = 0.05

if p_value < alpha:
    print('REJECT H0 - od_total_time and start_scan_to_end_scan are different')

else:
    print('FAIL TO REJECT H0 - od_total_time and start_scan_to_end_scan are similar')

    P-value : 0.7815123224221716
    FAIL TO REJECT H0 - od_total_time and start_scan_to_end_scan are similar
```

INFERENCE:

od_total_time and start_scan_to_end_scan are **similar**

Do hypothesis testing / visual analysis between actual_time aggregated value and OSRM time

- ✓ aggregated value (aggregated values are the values you'll get after merging the rows on the basis of trip_uuid)

```
df[['actual_time', 'osrm_time']].describe().T
```

✓ KDEPLOT:

```
plt.figure(figsize=(3,3))
sns.kdeplot(df['actual_time'])
sns.kdeplot(df['osrm_time'])
plt.show()
```

INFERENCE:

The **kdeplot** vividly shows that the graphs of both the groups are almost the **same distribution** and they have **almost same mean**. so ttest is performed.

✓ TTEST-IND:

```
# H0 - actual_time and osrm_time are similar
# Ha - actual_time and osrm_time are different

test_stat, p_value = spy.ttest_ind(df['actual_time'],df['osrm_time'])

print('P-value :',p_value)
alpha = 0.05

if p_value < alpha:
    print('REJECT H0 - actual_time and osrm_time are different')

else:
    print('FAIL TO REJECT H0 - actual_time and osrm_time are similar')

    P-value : 0.0
    REJECT H0 - actual_time and osrm_time are different
```

Visual Tests to know if the samples follow normal distribution

✓ HISTPLOT:

```
plt.figure(figsize = (10, 4))
sns.histplot(df['actual_time'], element = 'step', color = 'green')
sns.histplot(df['osrm_time'], element = 'step', color = 'lightblue')
plt.legend(['actual_time', 'osrm_time'])
plt.show()
```

✓ QQ Plot:

```
plt.figure(figsize = (15, 5))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for actual_time and osrm_time')
spy.probplot(df['actual_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for actual_time')
plt.subplot(1, 2, 2)
spy.probplot(df['osrm_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for osrm_time')
plt.show()
```


✓ Shapiro-Wilk test:

```
test_stat, p_value = shapiro(df['actual_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0
REJECT H0 - The sample does not follow normal distribution
```

```
test_stat, p_value = shapiro(df['osrm_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0
REJECT H0 - The sample does not follow normal distribution
```

✓ BOXCOX TRANSFORMATION:

```
transformed_actual_time = spy.boxcox(df['actual_time'])[0]
test_stat, p_value = spy.shapiro(transformed_actual_time)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 1.021792743086169e-28
REJECT H0 -The sample does not follow normal distribution
/usr/local/lib/python3.10/dist-packages/scipy/stats/_morestats.py:1882: UserWarning: p-value may not be accurate for N
    warnings.warn("p-value may not be accurate for N > 5000.")
```

```
transformed_osrm_time = spy.boxcox(df['osrm_time'])[0]
test_stat, p_value = spy.shapiro(transformed_osrm_time)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 -The sample follows normal distribution')

p-value 3.543600614978861e-35
REJECT H0 -The sample does not follow normal distribution
```

✓ Lavene's test:

```
# Null Hypothesis(H0) - Homogenous Variance

# Alternate Hypothesis(HA) - Non Homogenous Variance

test_stat, p_value = spy.levene(df['actual_time'], df['osrm_time'])
print('p-value', p_value)
if p_value < 0.05:
    print('The samples do not have Homogenous Variance')
else:
    print('The samples have Homogenous Variance ')

p-value 1.871297993683208e-220
The samples do not have Homogenous Variance
```

Since the samples do not follow any of the assumptions T-Test cannot be applied here, we can perform its non parametric equivalent test i.e., Mann-Whitney U rank test for two independent samples.

✓ Mann-Whitney U rank test:

```
test_stat, p_value = spy.mannwhitneyu(df['actual_time'], df['osrm_time'])
print('p-value', p_value)
alpha = 0.05
if p_value < alpha:
    print('REJECT H0 - actual_time and osrm_time are different')

else:
    print('FAIL TO REJECT H0 - actual_time and osrm_time are similar')

p-value 0.0
REJECT H0 - actual_time and osrm_time are different
```

INFERENCE:

actual_time and osrm_time are **different**

Do hypothesis testing/ visual analysis between actual_time aggregated value and segment actual time

- ✓ aggregated value (aggregated values are the values you'll get after merging the rows on the basis of trip_uuid)

```
df[['actual_time', 'segment_actual_time']].describe().T
```

- ✓ KDE PLOT:

```
plt.figure(figsize=(3,3))  
sns.kdeplot(df['actual_time'])  
sns.kdeplot(df['segment_actual_time'])  
plt.show()
```

INFERENCE:

The **kdeplot** vividly shows that the graphs of both the groups are the **same distribution** and they have **almost same mean**. so ttest is performed.

✓ TTEST:

```
# H0 - actual_time and segment_actual_time are similar
# Ha - actual_time and segment_actual_time are different

test_stat, p_value = spy.ttest_ind(df['actual_time'],df['segment_actual_time'])

print('P-value :',p_value)
alpha = 0.05

if p_value < alpha:
    print('REJECT H0 - actual_time and segment_actual_time are different')

else:
    print('FAIL TO REJECT H0 - actual_time and segment_actual_time are similar')

    P-value : 0.6165138648224772
    FAIL TO REJECT H0 - actual_time and segment_actual_time are similar
```

✓ QQ Plot:

```
plt.figure(figsize = (15, 5))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for actual_time and segment_actual_time')
spy.probplot(df['actual_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for actual_time')
plt.subplot(1, 2, 2)
spy.probplot(df['segment_actual_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for segment_actual_time')
plt.plot()
```

✓ Shapiro-Wilk test:

```
test_stat, p_value = spy.shapiro(df['actual_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 0.0
REJECT H0 - The sample does not follow normal distribution

test_stat, p_value = spy.shapiro(df['segment_actual_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0  
REJECT H0 - The sample does not follow normal distribution
```

✓ BOXCOX TRANSFORMATION:

```
transformed_actual_time = spy.boxcox(df['actual_time'])[0]  
test_stat, p_value = spy.shapiro(transformed_actual_time)  
print('p-value', p_value)  
if p_value < 0.05:  
    print('REJECT H0 -The sample does not follow normal distribution')  
else:  
    print('FAIL TO REJECT H0 - The sample follows normal distribution')  
  
p-value 1.021792743086169e-28  
REJECT H0 -The sample does not follow normal distribution  
/usr/local/lib/python3.10/dist-packages/scipy/stats/_morestats.py:1882: UserWarning: p-value may not be accurate for N  
    warnings.warn("p-value may not be accurate for N > 5000.")
```

```
transformed_segment_actual_time = spy.boxcox(df['segment_actual_time'])[0]  
test_stat, p_value = spy.shapiro(transformed_segment_actual_time)  
print('p-value', p_value)  
if p_value < 0.05:  
    print('REJECT H0 -The sample does not follow normal distribution')  
else:  
    print('FAIL TO REJECT H0 - The sample follows normal distribution')  
  
p-value 5.696120172016859e-29  
REJECT H0 -The sample does not follow normal distribution
```

✓ LAVENE'S test:


```
# Null Hypothesis(H0) - Homogenous Variance

# Alternate Hypothesis(HA) - Non Homogenous Variance

test_stat, p_value = spy.levene(df['actual_time'], df['segment_actual_time'])
print('p-value', p_value)
if p_value < 0.05:
    print('The samples do not have Homogenous Variance')
else:
    print('The samples have Homogenous Variance ')

p-value 0.6955022668700895
The samples have Homogenous Variance
```

Since the samples do not follow any of the assumptions T-Test cannot be applied here, we can perform its non parametric equivalent test i.e., Mann-Whitney U rank test for two independent samples.

✓ Mann-Whitney U rank test:

```
test_stat, p_value = spy.mannwhitneyu(df['actual_time'], df['segment_actual_time'])
print('p-value', p_value)
alpha = 0.05
if p_value < alpha:
    print('REJECT H0 - actual_time and segment_actual_time are different')

else:
    print('FAIL TO REJECT H0 - actual_time and segment_actual_time are similar')

p-value 0.4164235159622476
FAIL TO REJECT H0 - actual_time and segment_actual_time are similar
```

INFERENCE:

actual_time and segment_actual_time are **similar**

Do hypothesis testing/ visual analysis between osrm distance aggregated value and segment osrm

- ✓ distance aggregated value (aggregated values are the values you'll get after merging the rows on the basis of trip_uuid)

```
df[['osrm_distance', 'segment_osrm_distance']].describe().T
```

- ✓ KDE PLOT:

```
plt.figure(figsize=(3,3))  
sns.kdeplot(df['osrm_distance'])  
sns.kdeplot(df['segment_osrm_distance'])  
plt.show()
```

▼ TTEST:

```
# H0 - osrm_distance and segment_osrm_distance are similar
# Ha - osrm_distance and segment_osrm_distance are different

test_stat, p_value = spy.ttest_ind(df['osrm_distance'],df['segment_osrm_distance'])

print('P-value :',p_value)
alpha = 0.05

if p_value < alpha:
    print('REJECT H0 - osrm_distance and segment_osrm_distance are different')

else:
    print('FAIL TO REJECT H0 - osrm_distance and segment_osrm_distance are similar')

    P-value : 3.842631473353718e-05
    REJECT H0 - osrm_distance and segment_osrm_distance are different
```

✓ QQ Plot:

```
plt.figure(figsize = (15, 5))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for osrm_distance and segment_osrm_distance')
spy.probplot(df['osrm_distance'], plot = plt, dist = 'norm')
plt.title('QQ plot for osrm_distance')
plt.subplot(1, 2, 2)
spy.probplot(df['segment_osrm_distance'], plot = plt, dist = 'norm')
plt.title('QQ plot for segment_osrm_distance')
plt.plot()
```

✓ Shapiro-Wilk test:

```
test_stat, p_value = spy.shapiro(df['osrm_distance'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 0.0
REJECT H0 - The sample does not follow normal distribution

test_stat, p_value = spy.shapiro(df['segment_osrm_distance'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 0.0
REJECT H0 - The sample does not follow normal distribution
```

✓ BOXCOX TRANSFORMATION:

```
transformed_osrm_distance = spy.boxcox(df['osrm_distance'])[0]
test_stat, p_value = spy.shapiro(transformed_osrm_distance)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 7.114532433223529e-41
REJECT H0 -The sample does not follow normal distribution
```

```
transformed_segment_osrm_distance = spy.boxcox(df['segment_osrm_distance'])[0]
test_stat, p_value = spy.shapiro(transformed_segment_osrm_distance)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 3.0623432935550394e-38
REJECT H0 -The sample does not follow normal distribution
```

✓ LAVENE'S TEST:

```
# Null Hypothesis(H0) - Homogenous Variance

# Alternate Hypothesis(HA) - Non Homogenous Variance

test_stat, p_value = spy.levene(df['osrm_distance'], df['segment_osrm_distance'])
print('p-value', p_value)
if p_value < 0.05:
    print('The samples do not have Homogenous Variance')
else:
    print('The samples have Homogenous Variance ')

p-value 0.00020976354422600578
The samples do not have Homogenous Variance
```

✓ Mann-Whitney U rank test:

```
test_stat, p_value = spy.mannwhitneyu(df['osrm_distance'], df['segment_osrm_distance'])
print('p-value', p_value)
alpha = 0.05
if p_value < alpha:
    print('REJECT H0 - osrm_distance and segment_osrm_distance are different')

else:
    print('FAIL TO REJECT H0 - osrm_distance and segment_osrm_distance are similar')

p-value 9.511383588276375e-07
REJECT H0 - osrm_distance and segment_osrm_distance are different
```

INFERENCE:

osrm_distance and segment_osrm_distance are **different**

Do hypothesis testing/ visual analysis between osrm time aggregated value and segment osrm time

- ✓ aggregated value (aggregated values are the values you'll get after merging the rows on the basis of trip_uuid)

```
df[['osrm_time', 'segment_osrm_time']].describe().T
```

- ✓ KDE PLOT:

```
plt.figure(figsize=(3,3))
sns.kdeplot(df['osrm_time'])
sns.kdeplot(df['segment_osrm_time'])
plt.show()
```

✓ TTEST:

```
# H0 - osrm_time and segment_osrm_time are similar
# Ha - osrm_time and segment_osrm_time are different

test_stat, p_value = spy.ttest_ind(df['osrm_time'],df['segment_osrm_time'])

print('P-value :',p_value)
alpha = 0.05

if p_value < alpha:
    print('REJECT H0 - osrm_time and segment_osrm_time are different')

else:
    print('FAIL TO REJECT H0 - osrm_time and segment_osrm_time are similar')
```


P-value : 9.956426798219171e-09

REJECT H_0 - osrm_time and segment_osrm_time are different

✓ QQ Plot:

```
plt.figure(figsize = (15, 5))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for osrm_time and segment_osrm_time')
spy.probplot(df['osrm_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for osrm_time')
plt.subplot(1, 2, 2)
spy.probplot(df['segment_osrm_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for segment_osrm_time')
plt.plot()
```

✓ Shapiro-Wilk test:

```
test_stat, p_value = spy.shapiro(df['osrm_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0
REJECT H0 - The sample does not follow normal distribution
```

```
test_stat, p_value = spy.shapiro(df['segment_osrm_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 - The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 0.0
REJECT H0 - The sample does not follow normal distribution
```

✓ BOXCOX TRANSFORMATION:

```
transformed_osrm_time = spy.boxcox(df['osrm_time'])[0]
test_stat, p_value = spy.shapiro(transformed_osrm_time)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')
```

```
p-value 3.543600614978861e-35
REJECT H0 -The sample does not follow normal distribution
```

```
transformed_segment_osrm_time = spy.boxcox(df['segment_osrm_time'])[0]
test_stat, p_value = spy.shapiro(transformed_segment_osrm_time)
print('p-value', p_value)
if p_value < 0.05:
    print('REJECT H0 -The sample does not follow normal distribution')
else:
    print('FAIL TO REJECT H0 - The sample follows normal distribution')

p-value 4.893250997154572e-34
REJECT H0 -The sample does not follow normal distribution
```

✓ LAVENE'S TEST:

```
# Null Hypothesis(H0) - Homogenous Variance

# Alternate Hypothesis(HA) - Non Homogenous Variance

test_stat, p_value = spy.levene(df['osrm_time'], df['segment_osrm_time'])
print('p-value', p_value)
if p_value < 0.05:
    print('The samples do not have Homogenous Variance')
else:
    print('The samples have Homogenous Variance ')

p-value 8.349482669010088e-08
The samples do not have Homogenous Variance
```

✓ Mann-Whitney U rank test:

```
test_stat, p_value = spy.mannwhitneyu(df['osrm_time'], df['segment_osrm_time'])
print('p-value', p_value)
alpha = 0.05
if p_value < alpha:
    print('REJECT H0 - osrm_time and segment_osrm_time are different')

else:
    print('FAIL TO REJECT H0 - osrm_time and segment_osrm_time are similar')

    p-value 2.2995370859748865e-08
    REJECT H0 - osrm_time and segment_osrm_time are different
```

INFERENCE:

osrm_time and segment_osrm_time are **different**

- ✓ Find outliers in the numerical variables (you might find outliers in almost all the variables), and check it using visual analysis

```
numerical_columns = ['od_total_time', 'start_scan_to_end_scan', 'actual_distance_to_destination',
                     'actual_time', 'osrm_time', 'osrm_distance', 'segment_actual_time',
                     'segment_osrm_time', 'segment_osrm_distance']

df[numerical_columns].describe().T
```

▼ HISTPLOT:

```
plt.figure(figsize = (18,18))
for i in range(len(numerical_columns)):
    plt.subplot(3, 3, i + 1)
    sns.histplot(df[numerical_columns[i]], bins = 1000, kde = True)
    plt.title(f"Distribution of {numerical_columns[i]} column")
plt.show()
```

INFERENCE:

It can be inferred from the above plots that data in all the numerical columns are **right skewed**.

✓ BOX PLOT:

```
plt.figure(figsize = (18, 15))
for i in range(len(numerical_columns)):
    plt.subplot(3, 3, i + 1)
    sns.boxplot(df[numerical_columns[i]])
    plt.title(f"Distribution of {numerical_columns[i]} column")
plt.show()
```


INFERENCE:

It can be clearly seen in the above plots that there are outliers in all the numerical columns that needs to be treated.

✓ Detecting Outliers:

```
for i in numerical_columns:
    Q1 = np.quantile(df[i], 0.25)
    Q3 = np.quantile(df[i], 0.75)
    IQR = Q3 - Q1
    LB = Q1 - 1.5 * IQR
    UB = Q3 + 1.5 * IQR
    outliers = df.loc[(df[i] < LB) | (df[i] > UB)]
    print('Column :', i)
    print(f'Q1 : {Q1}')
    print(f'Q3 : {Q3}')
    print(f'IQR : {IQR}')
    print(f'LB : {LB}')
    print(f'UB : {UB}')
    print(f'Number of outliers : {outliers.shape[0]}')
    print('-----')
```

```
Column : od_total_time
Q1 : 149.93
Q3 : 638.2
IQR : 488.27000000000004
LB : -582.4750000000001
UB : 1370.605
Number of outliers : 1266
```

```
-----
Column : start_scan_to_end_scan
Q1 : 149.0
Q3 : 637.0
IQR : 488.0
LB : -583.0
UB : 1369.0
```

```
Number of outliers : 1267
-----
Column : actual_distance_to_destination
Q1 : 22.83723905859321
Q3 : 164.58320763841138
IQR : 141.74596857981817
LB : -189.78171381113404
UB : 377.2021605081386
Number of outliers : 1449
-----
Column : actual_time
Q1 : 67.0
Q3 : 370.0
IQR : 303.0
LB : -387.5
UB : 824.5
Number of outliers : 1643
-----
Column : osrm_time
Q1 : 29.0
Q3 : 168.0
IQR : 139.0
LB : -179.5
UB : 376.5
Number of outliers : 1517
-----
Column : osrm_distance
Q1 : 30.8192
Q3 : 208.475
IQR : 177.6558
LB : -235.6645
UB : 474.9587
Number of outliers : 1524
-----
Column : segment_actual_time
Q1 : 66.0
Q3 : 367.0
IQR : 301.0
LB : -385.5
UB : 818.5
Number of outliers : 1643
-----
```

Column • segment osrm time

INFERENCE:

The outliers present in our sample data can be the true outliers. It's best to remove outliers only when there is a sound reason for doing so. Some outliers represent natural variations in the population, and they should be left as is in the dataset.

✓ Do one-hot encoding of categorical variables

```
df.nunique()
```

```
data                2
trip_creation_time  14817
route_schedule_uuid 1504
route_type          2
trip_uuid           14817
source_center       1508
source_name         1499
destination_center  1481
destination_name    1469
od_start_time       26369
od_end_time         26369
start_scan_to_end_scan 1915
actual_distance_to_destination 144515
actual_time         3182
osrm_time           1531
osrm_distance       138046
segment_actual_time  747
segment_osrm_time   214
segment_osrm_distance 113799
dtype: int64
```

Columns with 2 unique values can be encoded. Here, its only data and route type columns.

▼ Route type

```
# Get value counts before one-hot encoding
```

```
df['route_type'].value_counts()
```

```
Carting    8908  
FTL        5909  
Name: route_type, dtype: int64
```

```
# Perform one-hot encoding on categorical column route type
```

```
label_encoder = LabelEncoder()  
df['route_type'] = label_encoder.fit_transform(df['route_type'])
```

```
# Get value counts after one-hot encoding
```

```
df['route_type'].value_counts()
```

```
0    8908  
1    5909  
Name: route_type, dtype: int64
```

▼ DATA TYPE:

```
# Get value counts of categorical variable 'data' before one-hot encoding
```

```
df['data'].value_counts()
```

```
training    10654
test        4163
Name: data, dtype: int64
```

```
# Perform one-hot encoding on categorical variable 'data'
```

```
label_encoder = LabelEncoder()
df['data'] = label_encoder.fit_transform(df['data'])
```

```
# Get value counts after one-hot encoding
```

```
df['data'].value_counts()
```

```
1    10654
0     4163
Name: data, dtype: int64
```

✓ Normalize/ Standardize the numerical features using MinMaxScaler or StandardScaler.

```
plt.figure(figsize = (6, 3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['od_total_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['od_total_time']} column")
plt.legend('od_total_time')
plt.plot()
```

```
plt.figure(figsize = (6,3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['start_scan_to_end_scan'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['start_scan_to_end_scan']} column")
plt.plot()
```

```
plt.figure(figsize = (6,3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['actual_distance_to_destination'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['actual_distance_to_destination']} column")
plt.plot()
```

```
plt.figure(figsize = (6,3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['actual_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['actual_time']} column")
plt.plot()
```



```
plt.figure(figsize = (6,3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['osrm_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['osrm_time']} column")
plt.plot()
```



```
plt.figure(figsize = (6,3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['osrm_distance'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['osrm_distance']} column")
plt.plot()
```

```
plt.figure(figsize = (6,3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['segment_actual_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['segment_actual_time']} column")
plt.plot()
```

```
plt.figure(figsize = (6,3))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df['segment_osrm_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"Normalized {df['segment_osrm_time']} column")
plt.plot()
```